

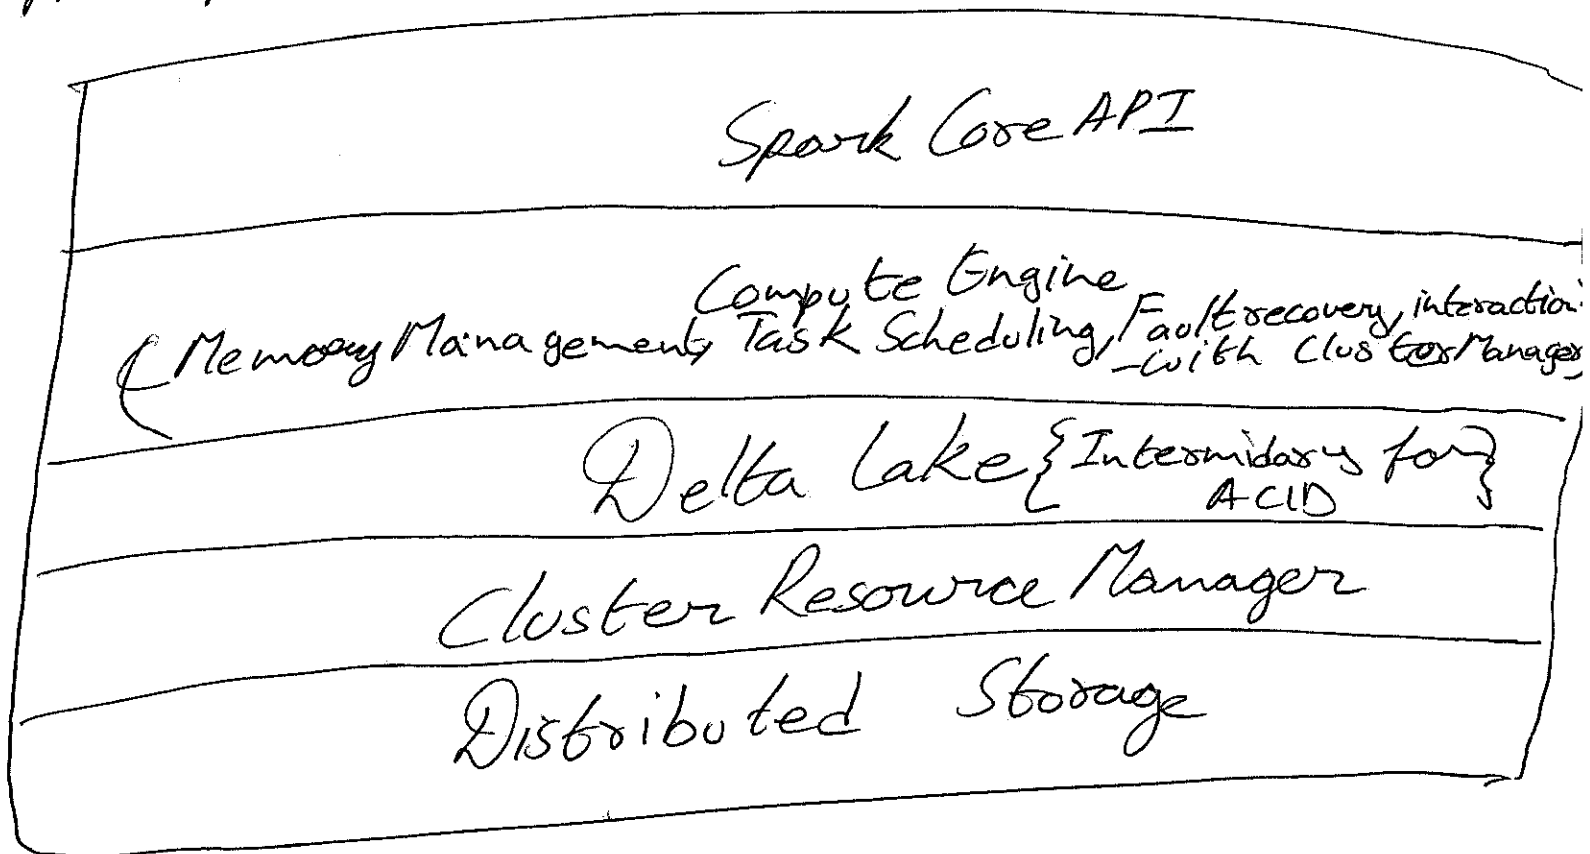
0

Delta lake for  
Apache Spark  
(How does it work?)

What is Delta Lake:-

①

Delta Lake provides an intermediary layer ~~for~~ in between Spark & Storage Layer to provide ACID compliance. As, storage layer in Hadoop is typically ~~the~~ the non ACID part.



Start Delta lake:-

```
$ spark-shell --master yarn \
```

```
--jars="delta/delta-core-2.11-0.3.0.jar"
```

(or)

```
$ spark-shell --master yarn \
```

```
--packages io.delta:delta-core-2.11:0.3.0
```

(2)

Code changes for Delta:-

Initial:-

```
→ spark.range(100).repartition(1).write.  
  .mode("overwrite").csv("/tmp/test-1/")
```

```
→ scala.util.Try(spark.range(100).  
  repartition(1).map { i =>  
    if (i > 50) {  
      Thread.sleep(5000)  
      throws new RuntimeException("Oops!")  
    }  
    i  
  }.write.mode("overwrite").csv("/tmp/test-2/"))
```

Changed:-

```
→ spark.range(100).repartition(1).write.mode(""  
  overwrite"").format("delta")  
  .save("/tmp/test-2/")
```

```
→ scala.util.Try(spark.range(100).repartition(1)  
  .map { i =>  
    if (i > 50) {
```

③

```
Thread.sleep(5000)
throw new RuntimeException("Ops!")
}
i
}
.write.mode("overwrite").format
("delta").save("/tmp/test-21")
```

but fails because of the mismatch in the column names, as shown below:-

```
> spark.range(100).printSchema
--id: long (nullable = false)
```

```
> spark.range(100).repartition(1).map(i =>
    if (i > 50) {
```

```
    Thread.sleep(5000)
    throw new RuntimeException("Ops!")
    }
    i
```

```
    }.printSchema
--value: long (nullable = true)
```

(4)

Rechanged ~~it~~ to fit the Column name —

```
> spark.range(100).repartition(1).map { i =>
  if (i > 50) {
    Thread.sleep(5000)
    throw new RuntimeException("Ops!")
  }
  i
} .select($"value".as("id"))
  .write.mode("overwrite")
  .format("delta")
  .save("/tmp/delta-test-1/")
```

Describe the Runtime error data is not  
lost.  
Testing for Data Atomicity —  
> spark.read.format("delta").load("/tmp/test-1/").count

test

(5)

Hence, Delta lake uses log files to in  
 "\_delta\_log" directory to provide Atomizity  
 for Runtime errors & to ~~read the~~ maintain  
 a list of Data additions & deletions.

Sample ETL:-

```
$ hadoop fs -cat /tmp/data-batch1/file
FName, LName, Phone, Age
Salman, Khan, 9876543210, 52
Akshay, Kumar, 9823456789, 48
```

## Back End Data loading:-

```
> val df = spark.read.format("csv").option
  ("header", "true").option
  ("inferSchema", "true").load
  ("tmp/data-batch1")
```

```
> val df1 = df.select($"FName", $"LName",
  $"Phone", $"Age",
  (when($"Age" > 50, "old").otherwise("Young"))
  .alias("AgeGroup"))
```

```
> df1.write.format("delta").mode("overwrite")
  .save("tmp/stars")
```

```
> df1.show()
```

FName	LName	Phone	Age	AgeGroup
Salman	Khan	9876543210	52	old
Akshay	Kumar	9823456789	48	Young

## ~~## Data Manipulation:~~

6

### ## Reading Data from Delta Lake:-

```
> val df = spark.read.format("delta").load("/tmp/Stars")
```

↓  
Reads as a data frame.

```
> import io.delta.tables._
```

```
> val dt = DeltaTable.forPath(spark, "/tmp/Stars")
```

↓  
Reads as a "io.delta.tables.DeltaTable"

## With Delta Table ⇒ additions, Updates & conversion to standard Dataframe is possible.

```
> dt.delete("FName == 'Salman'")
```

```
> dt.toDF.show
```

FName	LName	Phone	Age	AgeGroup
Akshay	Kumar	983456789	48	Young

```
> dt.updateExpr("FName == 'Akshay'", Map("Age" → "Age+5"))
```

(7)

## ## Upsert using Delta Lake

```
> import io.delta.tables._
```

```
> val dt = DeltaTable.forPath(spark, "/tmp/stars")
```

```
> dt.toDF.show
```

FName	LName	Phone	Age	Age Group
Salman	Khan	9876543210	52	Old
Akshay	Kumar	9823456789	48	Young

```
$ hadoop fs -cat /tmp/data-batch2/*
```

FName, LName, Phone, Age

Akshay, Kumar, 683256789, 52 ⇒ Update

Saif Ali, Khan, 9123456789, 55 ⇒ New

Sanyam, Dutt, 98433212345678 ⇒ New

```
> val df = spark.read.format("csv").option(
  ("header", "true")).schema("FName
  STRING, LName STRING, Phone STRING,
  Age DOUBLE").load("/tmp/data-batch2/
```

```
> val df1 = df.select($"FName", $"LName", $"Phone",
  $"Age", (when($"Age" > 50, "Old")
  .otherwise("Young")) alias("AgeGroup"))
```



(8)

→ df1.show

New  
batch  
Data

FName	(Name	Phone	Age	AgeGroup
Akshay	Kumar	653256789	57.0	old
Saif Ali	Khan	9123456789	55	old
Sanjay	Dott	9843342345	58	old

→ dt.as("stars")  
 .merge(df1.as("inputs"), "stars.FName  
 = inputs.FName")

.whenMatched()

.updateExpr(  
 Map(  
 "Name" → "inputs.Name",  
 "Phone" → "inputs.phone",  
 "Age" → "inputs.Age",  
 "AgeGroup" → "inputs.AgeGroup"  
 )  
 )

Upsert  
operation

.whenNotMatched  
 .insertAll  
 .execute();

(9)

> dt.toDF.Show // New data inserted & Updated

FName	LName	Phone	Age	Age Group
Saif Ali	Khanna	9123456789	55	old
Akshay	Kumar	683256789	51	old
Sanjay	Dutt	918433212345	58	old
Salman	Khan	9876543210	52	old

## Time Travel using Delta Lake

> dt.history.show(false) // to get the  
// commit history

// Using commit versions

> spark.read.format("delta").option("versionAsOf", 10).load("/tmp/stars").show

> spark.read.format("delta").option("versionAsOf", 5).load("/tmp/stars").show

// Time travel using commit time

// There are few specifics of reading Data via commit time.

// User can't read before the first commit time & after the last commit time. Trying to read beyond the first & last commit time will result in an exception due to lack of data snapshot in these time periods.

//  $\text{commit-time-of-version}_0 \leq \text{commit-time-of-version}_1$   
will result in `Data@version_0`

// to get the most recent time stamp Data load delta lake with-out time stamp.

`> spark.read.format("delta").option("timestampAsOf",  
"2019-09-25 15:50:00.000").load("/tmp/stars").show`

↓  
for Data before 2019-09-25 15:50:00.000 time

`> spark.read.format("delta").load("/tmp/stars").show`  
↓  
for most recent Data.

Summary:-

Atomicity:- [Either all or nothing]

Failed job Data in Delta Doesn't get updated in the logs.  
Due to logging mechanism in between old Data is not lost.

Consistency:- [Data is always in the valid state]  
till Data writing is completed by the current GTL, the previous version is available consistently to the users. Only once the transactional log is updated the data view is changed.

Isolation:- [an operation must be isolated]  
Delta tables are auto refresh & Data is always available. Hence Isolation is achieved.

Durability:- [once committed data is never lost]  
Delta table are stored are distributed storage system. Hence, data is always ~~cons~~ available due to partition Tolerance.

Schema enforcement:-  
~~Schema is checked~~ Delta Lake automatically validates that the schema of the Dataframe being written is compatible with the Schema of the table.

Time Travel:- Delta Lake also provides built-in data versioning for easy roll backs & reproducing reports.

Small file problems:-

Small file problems in Delta Lake can be handled by "compaction" using "Auto OPTIMIZE" (or) "repartition"

Ex:- `val df = spark.read.format("delta").load("/some/data")`  
 ↓  
 1000 files

`df.repartition(10).write.format("delta").mode("overwrite").save("/some/data")`  
 ↓  
 10 files  
 [1000 original previous data  
 10 current compacted]

User can use "vacuum" command to delete the old data files. like: `VACUUM Table name`  
 (or)  
`VACUUM delta.`/data/events``

Update & Deletion:-

uses "merge", "whenMatched", "UpdateExpr", "WhenNotMatched", "insertAll" & "execute()" command to facilitate data addition & update

Using Optimize:-

`OPTIMIZE delta.`/data/events``

(or)

`OPTIMIZE events`

(or)

`OPTIMIZE events where date >= '2017-01-01'`

## Multistage Data Skipping:-

"Z-ordering" is a technique to colocate related information in same set of files. This co-locality is automatically used for Data Skipping in-addition to "partitionby" like:-

```
df.write.format("delta").mode("overwrite").  
partitionBy("Origin").save("/tmp/flights_delta")
```

```
spark.sql("DROP TABLE IF EXISTS flights")
```

```
spark.sql("create table flights Using DELTA  
LOCATION '/tmp/flights_delta'")
```

```
spark.sql("OPTIMIZE flights ZORDER BY (columnname)")
```